

UAI 2011

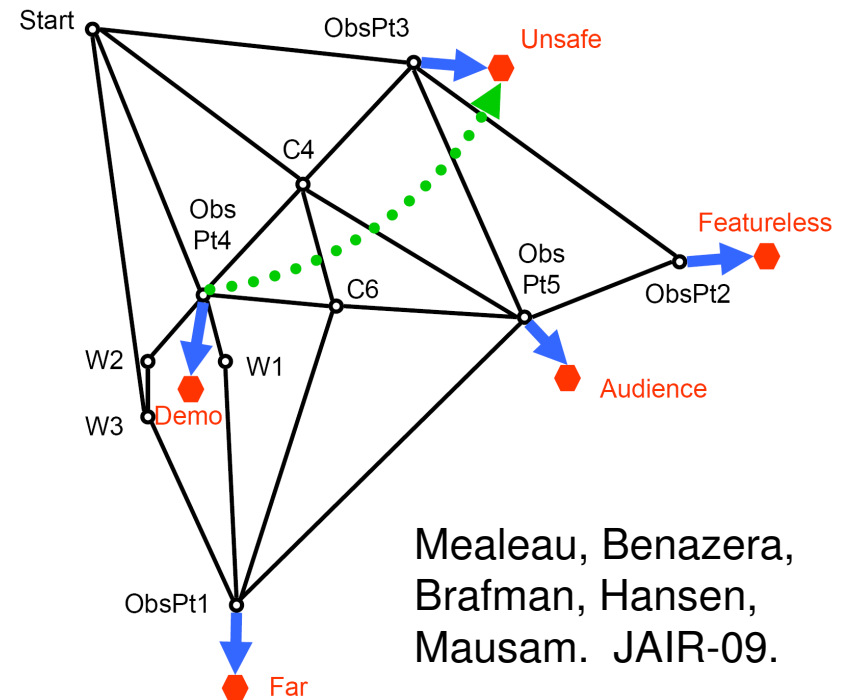

Symbolic Dynamic Programming for Discrete and Continuous State MDPs

Scott Sanner

Karina Valdivia Delgado
Leliane Nunes de Barros



Continuous Domains: Mars Rovers



Mealeau, Benazera,
Brafman, Hansen,
Mausam. JAIR-09.

- **Continuous state**
 - Time (t), Energy (e), Robot position (x, y, θ)
- **Closed-form exact solution?**
 - Currently only if 1D or piecewise rectilinear solution exists ☹

**Our work: exact for
multidimensional
nonlinear domains!**

Previous Work

Discrete and Continuous (DC-)MDPs

- Mixed discrete / continuous state

$$(\vec{b}, \vec{x}) = (b_1, \dots, b_n, x_1, \dots, x_m) \in \{0, 1\}^n \times \mathbb{R}^m$$

- Discrete action set $a \in \mathcal{A}$
- DBN factored transition model

$$P(\vec{b}', \vec{x}' | \vec{b}, \vec{x}, a) = \underbrace{\left(\prod_{i=1}^n P(b'_i | \vec{b}, \vec{x}, a) \right)}_{\text{discrete}} \underbrace{\left(\prod_{j=1}^m P(x'_j | \vec{b}, \vec{b}', \vec{x}, a) \right)}_{\text{continuous}}$$

- Arbitrary action-dependent reward

$$R_a(\vec{b}, \vec{x}) = x_1 + x_2$$

Value Iteration for DC-MDPs

- Value of policy in state is expected sum of rewards
- Want optimal value $V^{h,*}$ over horizons $h \in 0..H$
 - Implicitly provides optimal horizon-dependent policy
- Compute inductively via **Value Iteration** for $h \in 0..H$
 - **Regression step:**

$$Q_a^{h+1}(\vec{b}, \vec{x}) = R_a(\vec{b}, \vec{x}) + \gamma \cdot$$

$$\sum_{\vec{b}'} \int_{\vec{x}'} \left(\prod_{i=1}^n P(b'_i | \vec{b}, \vec{x}, a) \prod_{j=1}^m P(x'_j | \vec{b}, \vec{b}', \vec{x}, a) \right) V^h(\vec{b}', \vec{x}') d\vec{x}'$$

- **Maximization step:**

$$V_{h+1} = \max_{a \in A} Q_a^{h+1}(\vec{b}, \vec{x})$$

Exact Solutions to DC-MDPs: Domain

- 2-D Navigation

- State: $(x,y) \in \mathbb{R}^2$

- Actions:

- move-x-2

- $x' = x + 2$

- $y' = y$

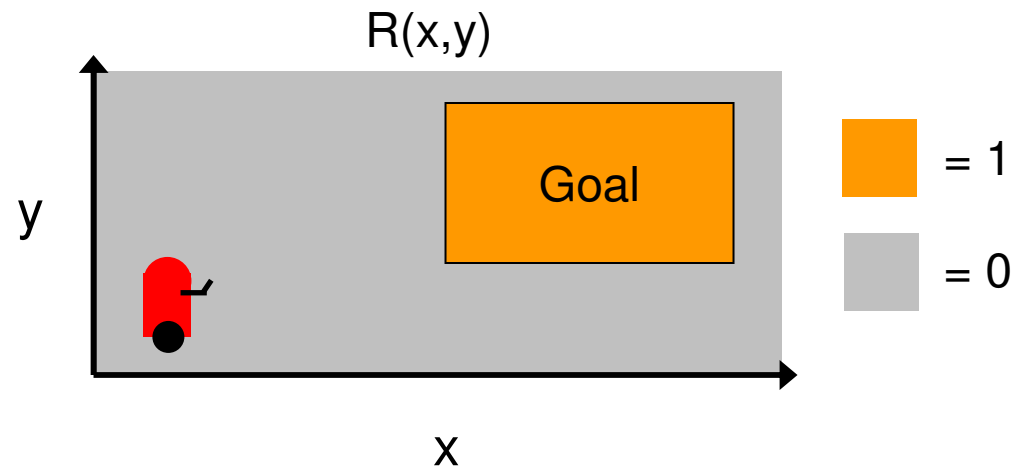
- move-y-2

- $x' = x$

- $y' = y + 2$

- Reward:

- $R(x,y) = \mathbb{I}[(x > 5) \wedge (x < 10) \wedge (y > 2) \wedge (y < 5)]$




Assumptions:

1. Continuous transitions are deterministic and linear
2. Discrete transitions can be stochastic
3. Reward is piecewise rectilinear

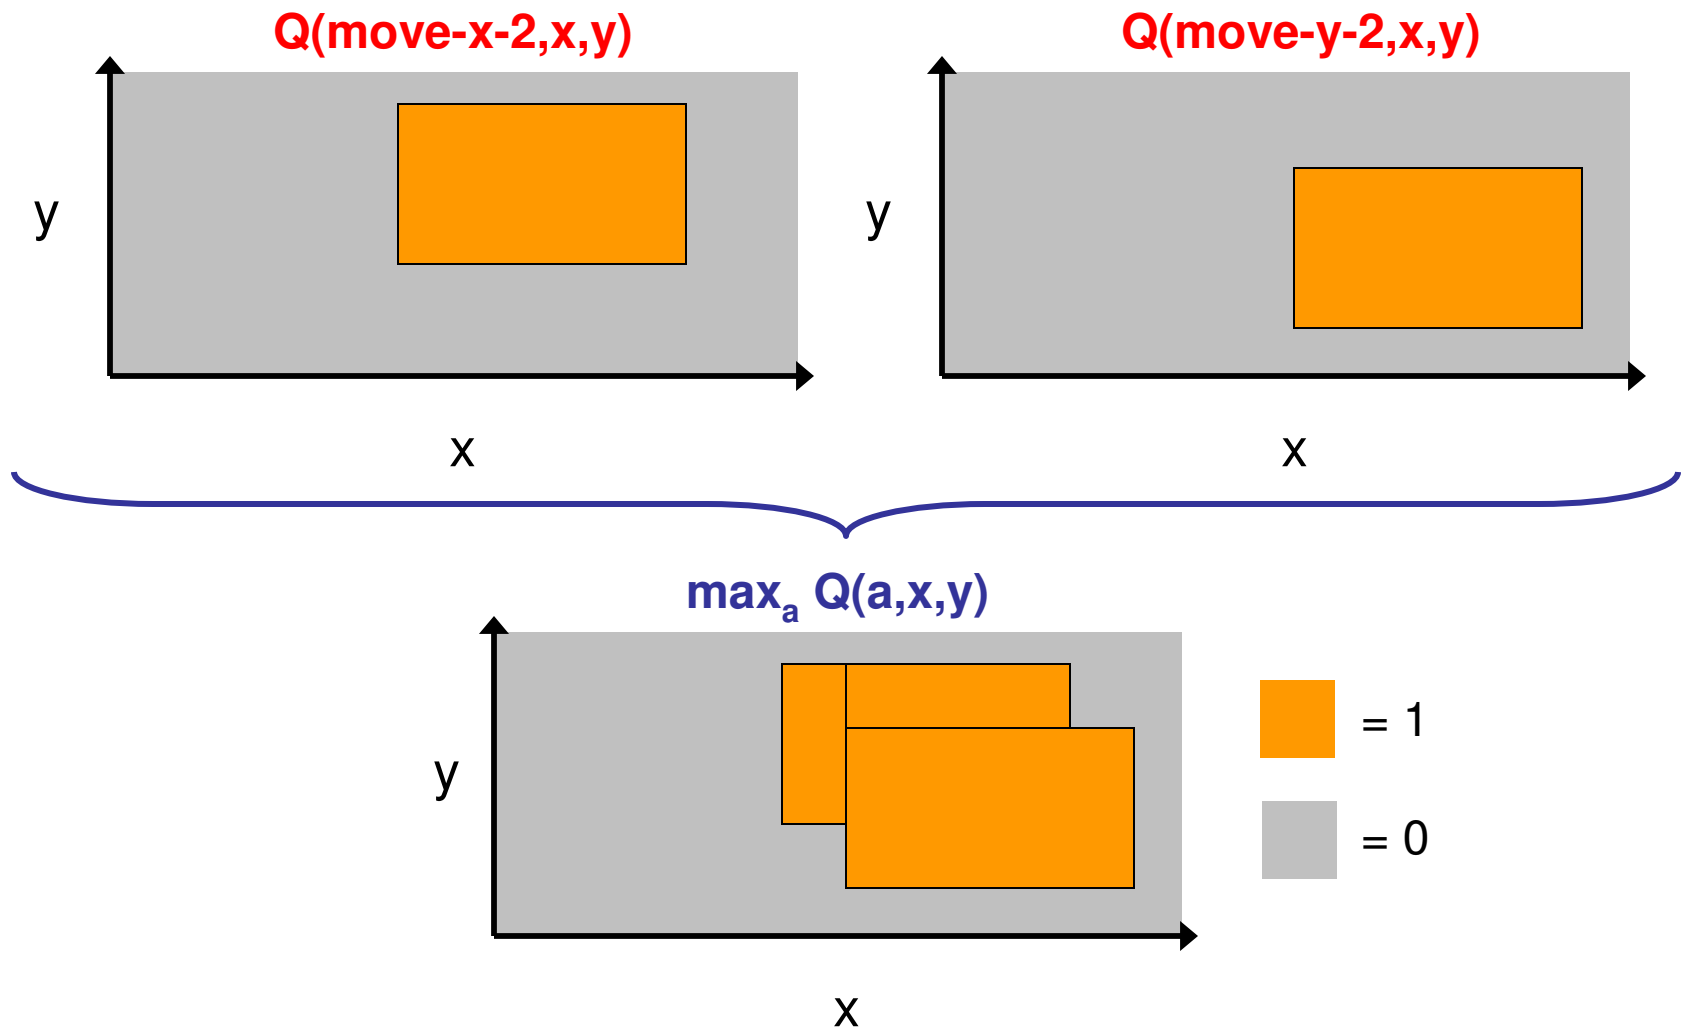
Exact Solutions to DC-MDPs: Regression

- Continuous regression is just translation of “pieces”



Exact Solutions to DC-MDPs: Maximization

- Q-value maximization yields piecewise rectilinear solution





Previous Work Limitations I

- Exact regression when transitions nonlinear?

Action **move-nonlin**:

$$- x' = x^3 y + y^2$$


$$- y' = y * \log(x^2 y)$$



How to compute
boundary in
closed-form?

Previous Work Limitations II

- $\max(\cdot, \cdot)$ when reward/value arbitrary piecewise?



A solution to previous limitations:

Symbolic Dynamic Programming (SDP)

n.b., motivated by SDP from
Boutilier *et al* (IJCAI-01) but here
continuous instead of relational

SDP uses Symbolic *Case* Representation

$$f = \begin{cases} \phi_1 : & f_1 \\ \vdots & \vdots \\ \phi_k : & f_k \end{cases}$$

$$P(x'|x, y) = \delta \left(x' - \begin{cases} (x < xy^2) \wedge (x > y) : & x + y \\ (x \geq xy^2) \vee (x \leq y) : & (x - y)^2 + 1 \end{cases} \right)$$

Deterministic transitions
represented by δ over
(conditional) equation

Logical combinations
of inequalities of
arbitrary expressions

Arbitrary
expressions

Case Operations: \oplus , \otimes

$$\left\{ \begin{array}{l} \phi_1 : f_1 \\ \phi_2 : f_2 \end{array} \right\} \oplus \left\{ \begin{array}{l} \psi_1 : g_1 \\ \psi_2 : g_2 \end{array} \right\} = \quad ?$$

Case Operations: \oplus , \otimes

$$\begin{cases} \phi_1 : f_1 \\ \phi_2 : f_2 \end{cases} \oplus \begin{cases} \psi_1 : g_1 \\ \psi_2 : g_2 \end{cases} = \begin{cases} \phi_1 \wedge \psi_1 : f_1 + g_1 \\ \phi_1 \wedge \psi_2 : f_1 + g_2 \\ \phi_2 \wedge \psi_1 : f_2 + g_1 \\ \phi_2 \wedge \psi_2 : f_2 + g_2 \end{cases}$$

- Similarly for \otimes
 - Expressions trivially closed under $+$, $*$
- What about \max ?
 - $\max(f_1, g_1)$ not pure arithmetic expression ☹

Case Operations: max

$$\max \left(\left\{ \phi_1 : f_1, \phi_2 : f_2 \right\}, \left\{ \psi_1 : g_1, \psi_2 : g_2 \right\} \right) = \quad ?$$

Case Operations: max

$$\max \left(\left\{ \phi_1 : f_1, \phi_2 : f_2 \right\}, \left\{ \psi_1 : g_1, \psi_2 : g_2 \right\} \right) = \begin{cases} \phi_1 \wedge \psi_1 \wedge f_1 > g_1 : f_1 \\ \phi_1 \wedge \psi_1 \wedge f_1 \leq g_1 : g_1 \\ \phi_1 \wedge \psi_2 \wedge f_1 > g_2 : f_1 \\ \phi_1 \wedge \psi_2 \wedge f_1 \leq g_2 : g_2 \\ \phi_2 \wedge \psi_1 \wedge f_2 > g_1 : f_2 \\ \phi_2 \wedge \psi_1 \wedge f_2 \leq g_1 : g_1 \\ \phi_2 \wedge \psi_2 \wedge f_2 > g_2 : f_2 \\ \phi_2 \wedge \psi_2 \wedge f_2 \leq g_2 : g_2 \end{cases}$$

**Key point: still in
case form!**

**Size blowup?
We'll get to that...**

Symbolic Dynamic Programming

- In a nutshell
 - $R(\cdot)$, $P(\cdot|\cdot)$ defined as case statements
 - Value iteration uses case operations
 - \oplus , \otimes , \max
 - Then provably:
 - $V^h(\cdot)$ is also in case form for all horizons h !
- Only “tricky part” is continuous regression

SDP Regression Step

- **Binary variables** b_i

- Factored $\sum_{b_i \in \{0,1\}}$ (e.g, SPUDD: Hoey *et al*, UAI-99)

- **Continuous variables** x_j

- $\int_x \delta[x - y] f(x) dx = f(y)$ triggers symbolic *substitution*, so

$$\int_{x'_j} \delta[x'_j - g(\vec{x})] V' dx'_j = V' \{x'_j / g(\vec{x})\}$$

- e.g.,

$$\int_{x'_1} \delta[x'_1 - (x_1^2 + 1)] \left(\begin{cases} \underline{x'_1} < 2 : & \underline{x'_1} \\ \underline{x'_1} \geq 2 : & \underline{x'^2_1} \end{cases} \right) dx'_1 = \begin{cases} \underline{x_1^2 + 1} < 2 : & \underline{x_1^2 + 1} \\ \underline{x_1^2 + 1} \geq 2 : & \underline{(x_1^2 + 1)^2} \end{cases}$$

- If g is *case*: need *conditional substitution*, see paper

Case \rightarrow XADD



SDP needs an efficient data structure for

- compact, minimal case representation
- efficient case operations

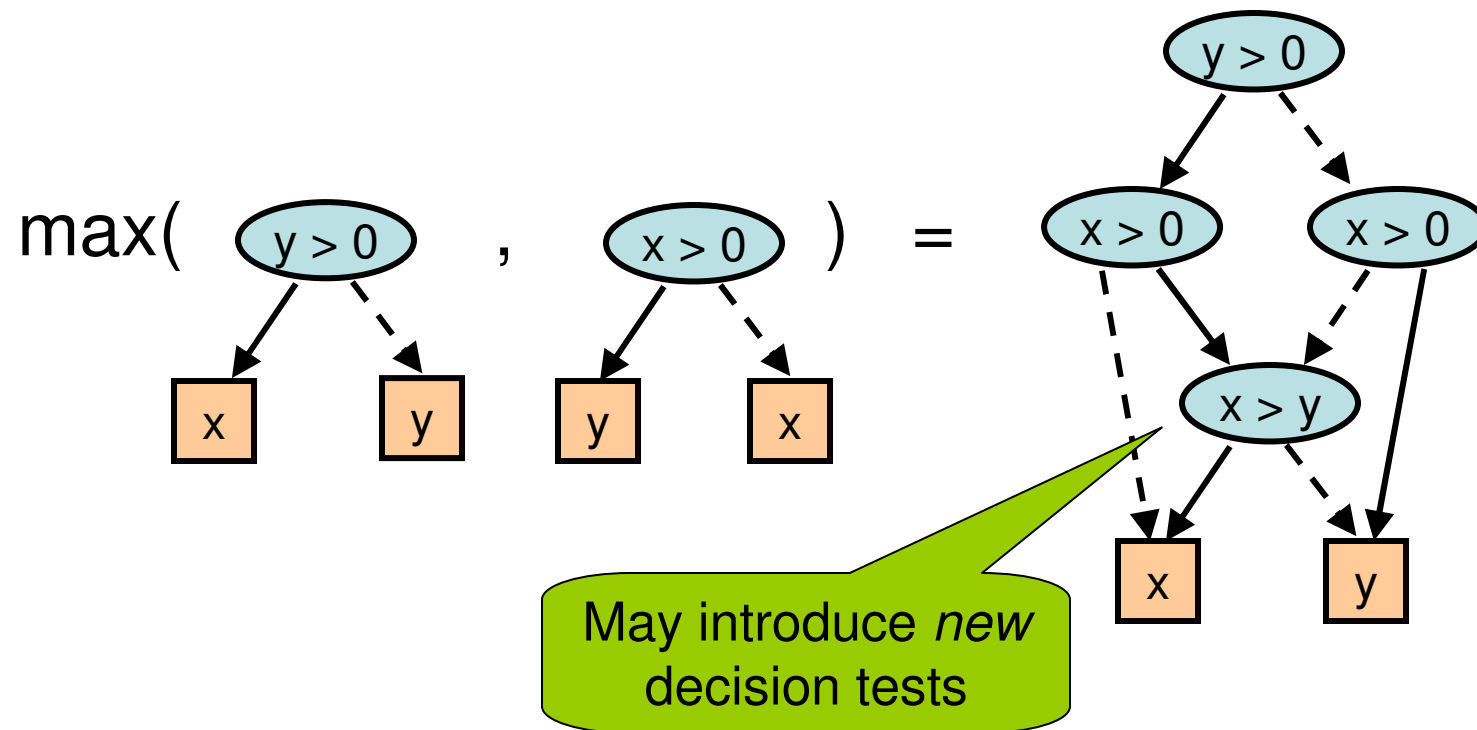
XADDs

- Extended ADD representation of case statements

$$V = \begin{cases} x_1 + k > 100 \wedge x_2 + k > 100 : & 0 \\ x_1 + k > 100 \wedge x_2 + k \leq 100 : & x_2 \\ x_1 + k \leq 100 \wedge x_2 + k > 100 : & x_1 \\ x_1 + x_2 + k > 100 \wedge x_1 + k \leq 100 \wedge x_2 + k \leq 100 \wedge x_2 > x_1 : & x_2 \\ x_1 + x_2 + k > 100 \wedge x_1 + k \leq 100 \wedge x_2 + k \leq 100 \wedge x_2 \leq x_1 : & x_1 \\ x_1 + x_2 + k \leq 100 : & x_1 + x_2 \end{cases}$$



XADD Maximization




Maintaining XADD Orderings I

- Max may get variables out of order


Decision
ordering
(root→leaf)

- $x > y$
• $y > 0$
• $x > 0$


$\max($



,



$=$




Newly introduced
node is out of order!

Maintaining XADD Orderings II

- Substitution may get vars out of order

Decision
ordering
(root→leaf):


- ↓
- $x > y$
 - $y > 0$
 - $x > z$



Correcting XADD Ordering


- Build *ordered* XADD from *unordered* XADD

z is out of order




Inductively assume ID₁
and ID₀ are ordered.

result will have z in order!



All operands ordered, so
applying \otimes , \oplus produces
ordered result!

XADD Pruning



Node unreachable –
 $x + y < 0$ always
false if $x > 0$ & $y > 0$

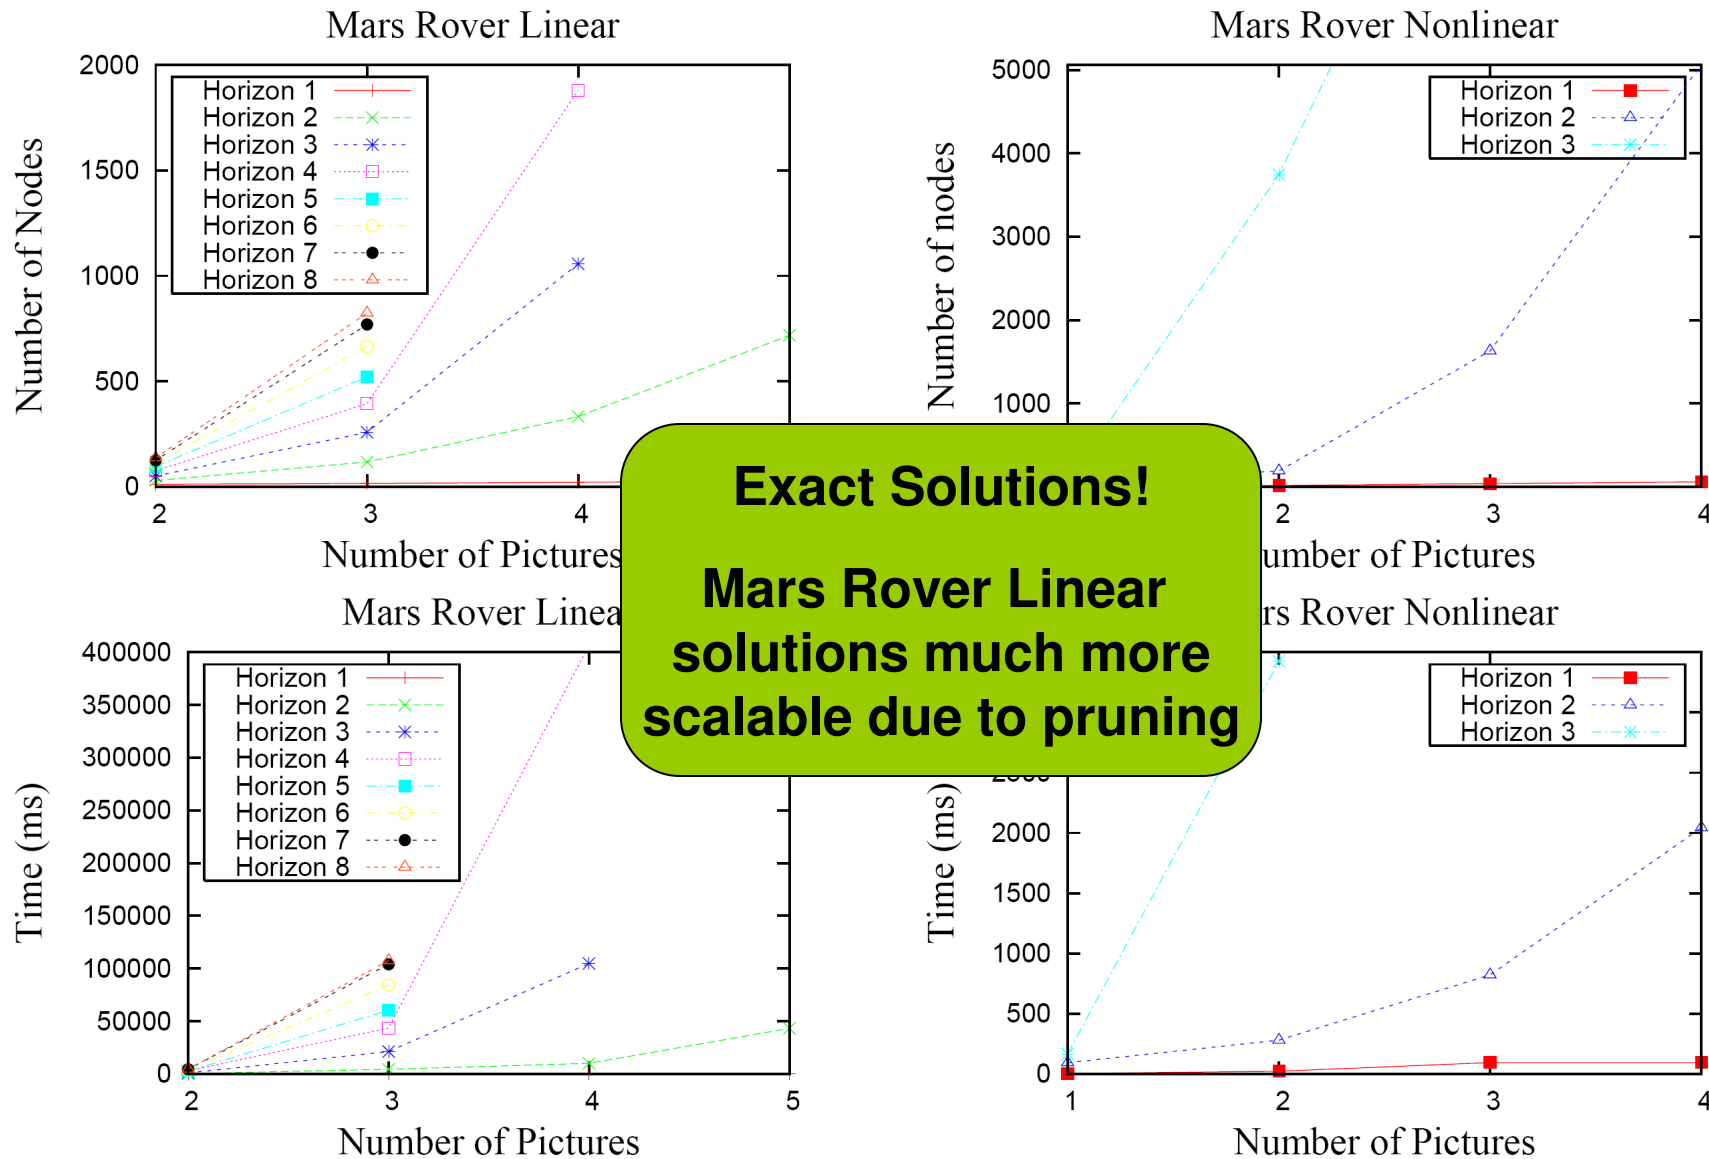
If **linear**, can detect
with feasibility checker
of LP solver & prune

Empirical Results



Problem Domains

- **Knapsack problem** (high-dimensional toy problem)
 - Transfer continuous resources to knapsack
 - Subject to capacity constraints
 - Reward for amount transferred
 - Can solve for optimal ∞ -horizon solution
- **Mars Rover** (variants of Bresina *et al*, UAI-02)
 - Linear
 - take pictures with linear time/energy constraints
 - Nonlinear
 - move to target (x,y) position, taking pictures along way
 - reward is truncated quadratic
- All problem domains / code online:
 - <http://code.google.com/p/xadd-inference/>

Results: Time and Space for Mars Rover



Results: XADD Pruning vs. No Pruning




Summary:

- without pruning: superlinear vs. horizon
- with pruning: linear vs. horizon


Worth the effort to prune!

Obligatory 3D Value Function Gallery


Knapsack



Mars Rover Linear



Mars Rover Nonlinear



Exact value functions in case form:

- **linear & nonlinear piecewise boundaries!**
- **nonlinear function surfaces!**

Conclusions

- First exact, closed-form solutions to subset of **multidimensional, nonlinear DC-MDPs**
- Key insights
 - Symbolic *case* representation
 - DP in terms of case \otimes , \oplus , \max
 - $\int \delta$ triggers (conditional) substitution
- Need compact case, efficient operations
 - Case \rightarrow Extended ADD (XADD)
 - \otimes , \oplus technique for efficient decision reordering
 - Advantages of pruning

Future Work

- Efficiency
 - XADD pruning in nonlinear case
 - XADD Approximation?
 - Extend APRICODD (St-Aubin *et al*, NIPS-00)
- Expressivity
 - Full continuous stochastic extension
 - Currently, continuous transitions are mixture of δ 's
 - Ideally want Gaussian noise, etc.
 - Continuous actions?
 - Partial observability?

Thank you!

Questions?

Extra Slides

XADD: Details

- The XADD is an ADD allowing
 - Arbitrary expressions at leaves
 - Arbitrary expression inequalities at decision nodes
 - If expressions polynomial, decisions & leaves have canonical form
 - Enforce ordering on all decision tests: $(x < y)$ before $(x < 3)$
- Operations same as XADD
 - But leaf operations may produce XADDs themselves!
 - May also require introduction of new decisions
 - E.g., maximization
- Can introduce support for substitution
 - Needed for SDP regression

1D: Boyan and Littman (1999)


- Exact Solutions to Time-dependent MDPs
 - Assume actions / transitions as follows

Action = bus:

$$t' = t + 30$$

$s' = \text{office}$

$$R(s,t) = -2 - t + 20 * I[s' = \text{office}]$$




Action = taxi:

$$t' = t + 10$$

$s' = \text{office}$


$$R(s,t) = -15 - t + 20 * I[s' = \text{office}]$$




1D: Boyan and Littman (1999)

- Continuous transitions are δ -functions
 - Regressions just sums & translations of value

$Q(a=\text{bus}, s=\text{office}, t)$





$Q(a=\text{taxi}, s=\text{office}, t)$




1D: Boyan and Littman (1999)

- Value max is just piecewise partitioning

$$\max(Q(a=\text{bus}, s=\text{office}, t), Q(a=\text{taxi}, s=\text{office}, t))$$



$$V(s=\text{office}, t)$$



Optimal solution is
piecewise linear