



# Relational and First-Order Decision-Theoretic Planning:

*Foundations and Future Directions*

Scott Sanner

`ssanner@cs.toronto.edu`

*Depth-Oral Presentation*


Department of Computer Science

University of Toronto

# Importance of DT Planning

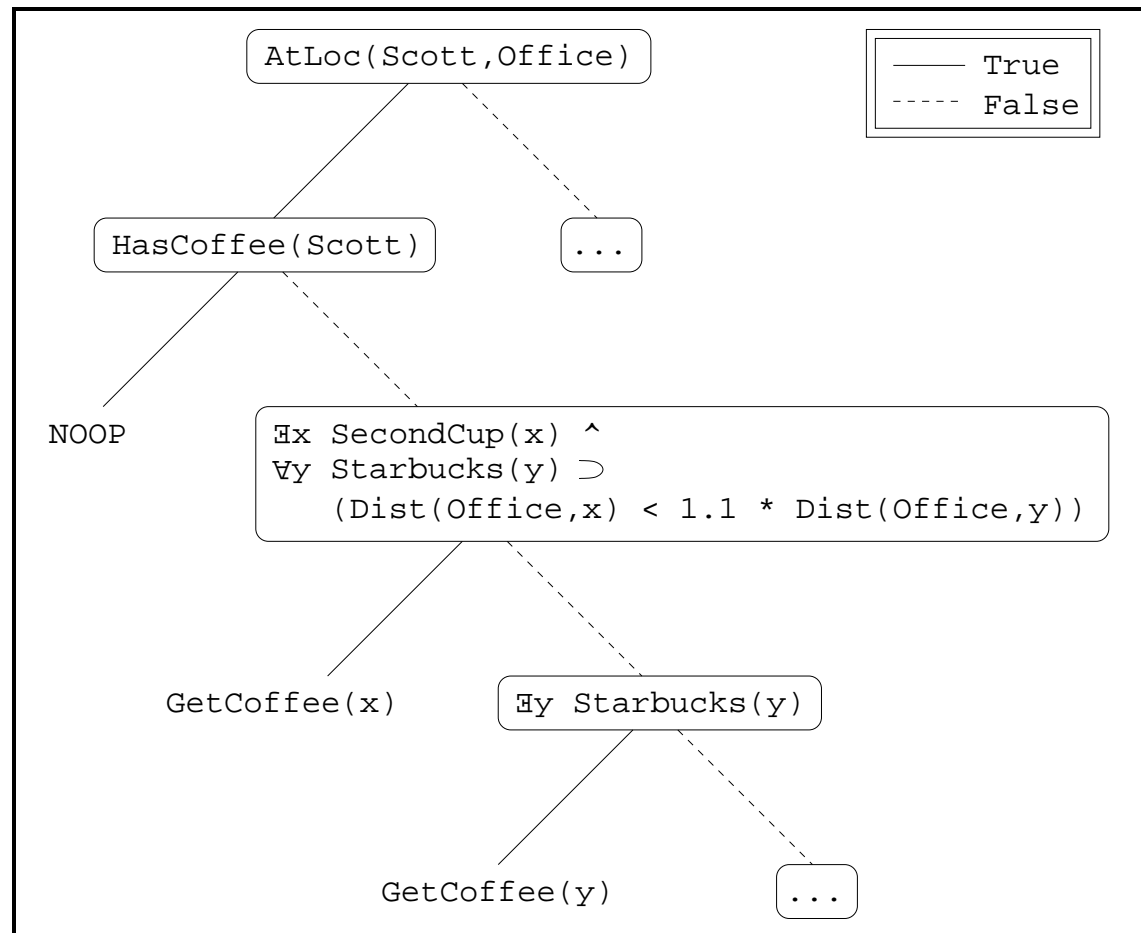


## Example: Planning the walk to the office.

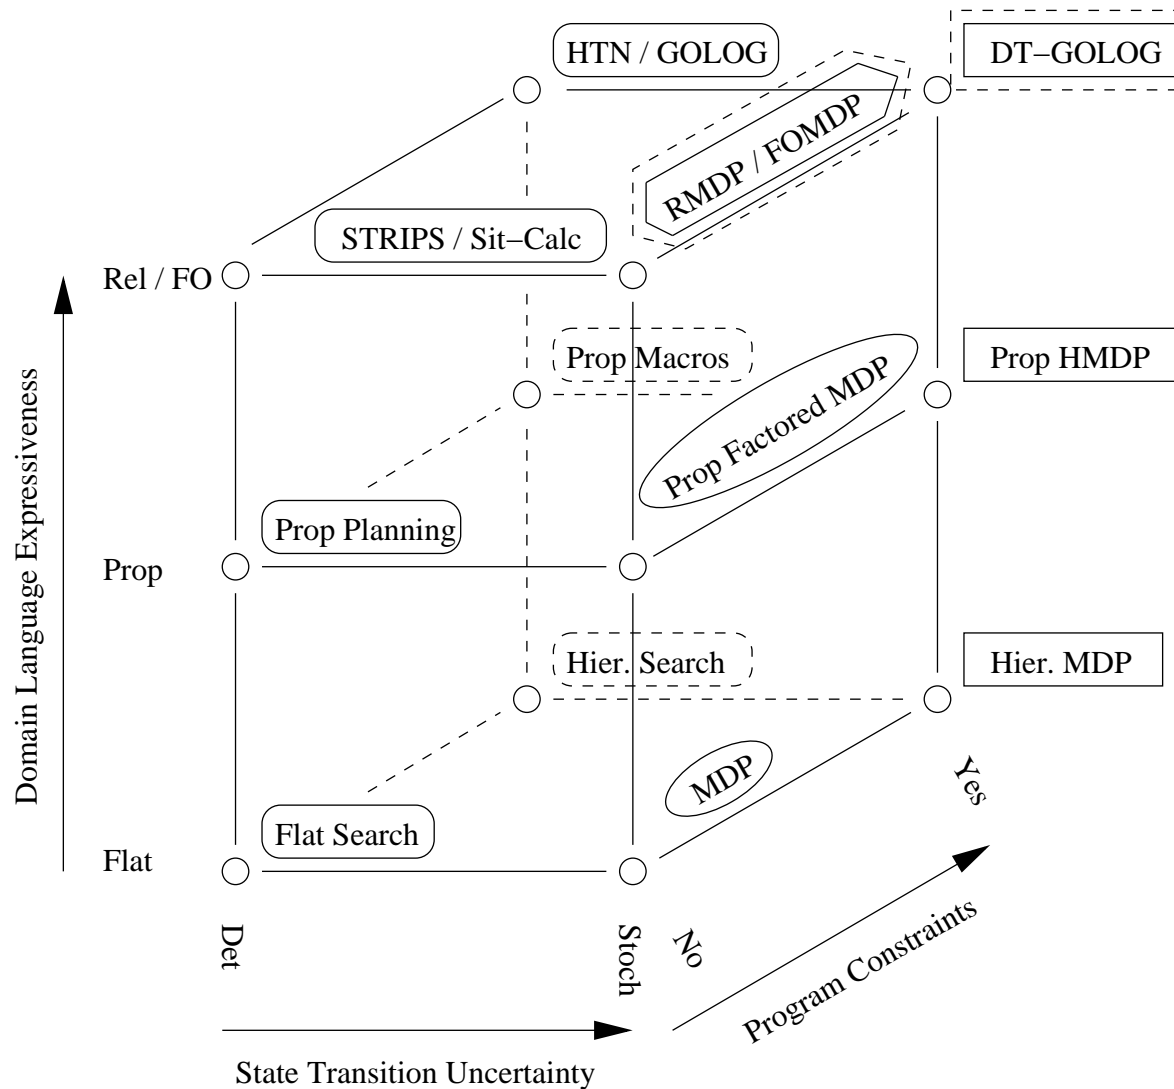


# Importance of FO/Rel. Representation

But how would we plan for all possible worlds?



# Presentation Overview



## Foundations:

Deterministic planning

MDPs and structured representations

Relational and first-order MDPs

Hierarchy and program constraints

## Future directions:

Structured representations in FOMDPs

Approximation in FOMDPs

FOMDPs and program constraints

# Det. Planning: Relational and FO

## STRIPS and PDDL

States:  $\{On(b_1, b_2), On(b_2, b_3)\}$

Actions:

```
(:action stack
:parameters (?a ?b)
:precondition (and (forall (?c) (not (on ?c ?a)) ... ))
:effect      (and (forall (?c)
                  (when (on ?a ?c)
                      (not (on ?a ?c))))
                (on ?a ?b)))
```

Problem: Dom.:  $b_1, b_2$ ; Init. state; Goal:  $\{On(x, y) \wedge On(y, z)\}$

## Situation calculus

Actions:  $stack(b_2, b_1)$ , Situations:  $s_0, do(stack(b_2, b_1), s_0)$ ,

Fluents:  $On(b_2, b_1, s_0)$ , SSAs:  $F(\vec{x}, do(a, s)) \equiv \Phi_F(\vec{x}, a, s)$

Problem: UNA, SSAs, action preconditions, initial situation

# Det. Planning: Program constraints

## GOLOG (LRLLS, 1997)

- Program  $\delta$  restricts execution:  $[a; b; [c|d]]^*$
- Specify execution using  $Do(\delta, s, s')$
- $Do(\delta, s, s')$  defined inductively on first argument:
  - *Primitive actions (base case)*
  - *Test actions*
  - *Sequence of programs*
  - *Nondeterministic {program choice, choice of arguments, iteration}*
  - *Other constructs: {if / then, while, procedures, programs}*
- Goal-Oriented Planning:  $Axioms \models (\exists s). Do(\delta, S_0, s) \wedge G(s)$

# MDP Representation (Informal)

- So far, we have discussed
  - Goal-oriented,
  - Relational / first-order representations, and
  - Program constraints.
- But how do we plan with
  - General reward functions / action costs?
  - Uncertain state transitions?
- One possible solution
  - Markov decision process (MDP) framework,
  - Assume infinite horizon,
  - Maximize expected sum of discounted future rewards.

# MDP Representation (Formal)

## ● MDP formally defined as $\langle S, A, T, R \rangle$

- $S$ : finite set of states
- $A$ : finite set of actions
- $T$ : transition function ( $T : S \times A \times S \rightarrow [0, 1]$ )
- $R$ : reward function ( $R : S \times A \rightarrow \mathbb{R}$ )

## ● Additionally define

- $\pi$ : policy mapping from states to actions ( $\pi : S \rightarrow A$ )
- $r^t$ : reward at time step  $t$
- $\gamma$ : discount factor where  $0 \leq \gamma < 1$
- $V_\pi(s) = E_\pi[\sum_{t=0}^{\infty} \gamma^t \cdot r^t | s_0 = s]$ : value of  $\pi$  starting from  $s$

## ● Goal: find optimal policy $\pi^*$


- $V^*(s) = V_{\pi^*}(s)$  maximizes value over all states



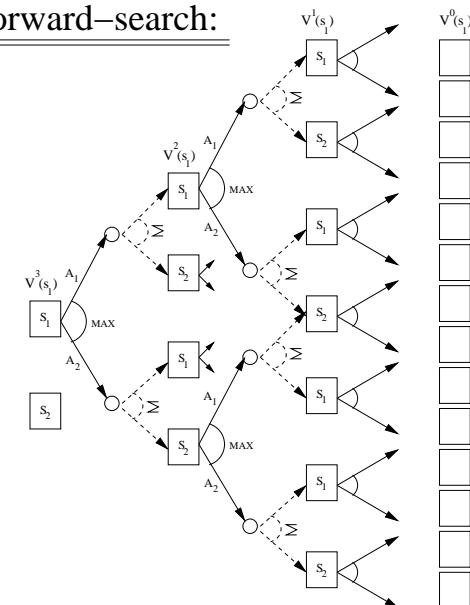
# MDP Solution Algorithms

- Dynamic programming
  - Value iteration
  - Policy iteration
  - Modified policy iteration
- Forward-search
- Real-time dynamic programming
- Linear program

Value iteration:



Forward-search:



# Prop. MDPs and Structured Repr. I

## Factored representation

- Transition DBN
- Influence diagram for reward


## Context-specific independence

- Transition CPT
- Reward structure


## Algorithms that exploit structure

- SPI (BDG,1995)
- SPUDD (HSHB,1999)
- APRICODD (SHB,2000)


Transition DBN:



ADD for  $P(S_3' | S_2, S_3, S_4)$ :



ADD for reward:



# Prop. MDPs and Structured Repr. II

- Exploit **additive structure** in value function
- Use **linear combination** of basis vectors:

$$V = w_1 a_1 + w_2 a_2 + w_3 a_3 \quad A = \{a_1, a_2, a_3\}$$

- Solve with **LP** or **approximate PI (API)**:

$$\vec{w}^{(t)} = \arg \min_{\vec{w}} \|(A - \gamma T_{\pi^{(t)}} A) \vec{w} - R_{\pi^{(t)}}\|$$

$$\pi^{(t+1)} = \arg \max_{\pi \in \Pi} (R_{\pi} + \gamma T_{\pi} A \vec{w}^{(t)})$$

- (GKP,2001) give **compact LP** for  $\mathcal{L}_{\infty}$  proj.
- **API  $\mathcal{L}_2$** : (KP,1999), **API  $\mathcal{L}_{\infty}$** : (GKP,2001), **LP  $\mathcal{L}_{\infty}$** : (SP,2001)

# Relational MDP Representation

## • PSTRIPS, PPDDL

```
(:action stack
:parameters (?a ?b)
:precondition (and (forall (?c) (and (not (on ?c ?a) ... ))))
:effect (probabilistic 0.7 (and (forall (?c)
                                (when (on ?a ?c)
                                    (not (on ?a ?c))))
                                (on ?a ?b))))

(define
  (:objects block0 block1 block2 ... )
  (:goal (exists (?b1 ?b2) (and (red ?b1) (on ?b1 ?b2)))))
```

- Concise representation of DT planning problem
- But action and state space can be extremely large even for small domain instantiations

# Relational MDP Algorithms

## First-order policy induction (YFG,2002; FYG,2003)

- Algorithm: Perform approximate policy iteration (using decision-list policies)

- Draw training samples  $D^{(t)}$  under a current policy  $\pi^{(t)}$
- Induce a new policy  $\pi^{(t+1)}$  from  $D^{(t)}$

- Example decision-list policy:

$$\begin{aligned} \text{goal-on}(a, b) \wedge \neg \text{on}(a, b) \wedge \text{holding}(a) &\longrightarrow \text{putdown}(a, b) \\ \text{goal-on}(a, b) \wedge \neg \text{on}(a, b) &\longrightarrow \text{pickup}(a) \end{aligned}$$

## Basis value function approximation (GKGK,2003)

- Algorithm: Compute basis value functions (using LP over sampled, weighted domains)
- Example value function:  $V(s) = V_{fo}^{otman}(f_1, e_1) + V_{fo}^{otman}(f_2, e_2)$

# First-order MDPs

## Representation

- Nature's choice for stochastic actions

$$P(\text{Putdown}S(b_1, b_2) \mid \text{Putdown}(b_1, b_2), s) = [\text{Wet}(b_1, s) : 0.7 ; \neg \text{Wet}(b_1, s) : 0.9]$$

$$P(\text{Putdown}F(b_1, b_2) \mid \text{Putdown}(b_1, b_2), s) = [\text{Wet}(b_1, s) : 0.3 ; \neg \text{Wet}(b_1, s) : 0.1]$$

- Reward and value function case partitions

$$\begin{aligned} \text{case} \quad [ & (\exists b_1, b_2). \text{On}(b_1, b_2, s) \wedge \text{Red}(b_1) : 10 ; \\ & \neg((\exists b_1, b_2). \text{On}(b_1, b_2, s) \wedge \text{Red}(b_1)) \wedge (\exists b_1, b_2). \text{On}(b_1, b_2, s) : 5 ; \\ & \neg(\exists b_1, b_2). \text{On}(b_1, b_2, s) : 0 ] \end{aligned}$$

## Solution approach


- Symbolic version of value iteration (BRP, 2001)
- Final  $\epsilon$ -optimal value function is a case partition

# MDPs with Program Constraints

## Approaches:

- Markov task decomposition  
(MHKPKDB, 1998)
- Macro-actions  
(SPS, 1999)
- MDP decomposition / abstraction  
(HMKDB, 1998)
- Program constraints  
(PR, 1998; AR, 2001/2)
- Macro-actions *and* program constraints  
(Diet, 1998)


## MDP Decomposition / Abstraction:




## Program Constraints:

Program  $P$  :


$[ [l \mid r] ; [u \mid d] ]^*$



MDP  $M$  :



$P \circ M$  :



# FOMDPs with Program Constraints

## DT-GOLOG representation (BRST,2000)

- Start with GOLOG's  $Do(\delta, s, s')$
- Generalize to a function representing the decision-theoretically optimal execution of  $\delta$ :  
 $BestDo(\delta, s) \rightarrow \mathbb{R}$
- DT nondeterministic program choice:*  
 $BestDo([\delta_1 | \delta_2]) = \max \{ BestDo(\delta_1) ; BestDo(\delta_2) \}$
- DT nondeterministic choice of arguments:*  
 $BestDo((\pi x)\delta(x)) = \max_x \{ BestDo(\delta(x)) \}$

## Solution approaches

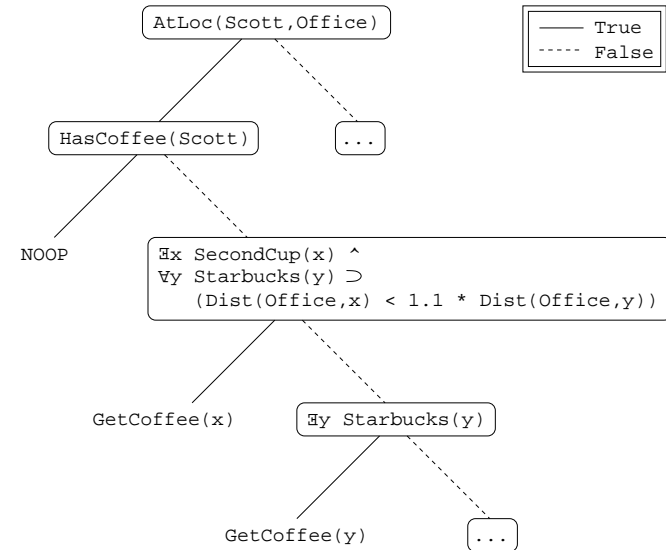
- Forward-search (BRST,2000)
- Forward-search with macros (FFL,2003)
- Real-time dynamic programming (Sout,2001)



# Future Directions I

## FOMDPs

- First-order ADDs and approximation
  - Partitions often have redundant structure
  - Exploit structure for computational and representational efficiency
  - Use APRICODD-style pruning for approximation
- First-order basis functions and approximation
  - Value function can be approximately additive
  - Exploit structure with weighted FO-basis functions



$$\begin{aligned}
 V(s) = & \\
 w_1 \cdot \text{case} & \left[ \begin{array}{l} (\exists c). \neg R(c, s) \wedge S(c) : 1 ; \\ (\forall c). S(c) \supset R(c, s) : 0 \end{array} \right] \\
 & + \\
 w_2 \cdot \text{case} & \left[ \begin{array}{l} (\exists c). \neg R(c, s) \wedge T(c) : 1 ; \\ (\forall c). T(c) \supset R(c, s) : 0 \end{array} \right]
 \end{aligned}$$

# Future Directions II

## FOMDPs

- DBN factored action decomposition
- Domain constraints
- First-order counting quantifiers / aggregators

$$\text{Reward} = \text{case}[(\exists_n b). \text{Red}(b) \wedge \text{In}(b, P, s) : 10 ; \\ \neg(\exists_n b). \text{Red}(b) \wedge \text{In}(b, P, s) : 0 ]$$

$$P(\text{Running}(c, s) \mid \text{NoReboot}(c, s) = \\ \frac{\#_d. \text{Connected}(c, d) \wedge \text{Working}(d)}{\#_d. \text{Connected}(c, d)}$$

## FOMDPs with program constraints

- First-order DT regression under program constraints
- Allows FODTR under prog. constraints without initial state knowledge

