


# **First-order Decision-theoretic Planning in Structured Relational Environments**

**Scott Sanner**  
**University of Toronto**  
*[ssanner@cs.toronto.edu](mailto:ssanner@cs.toronto.edu)*

# Why First-order DT Planning?

- Relational planning problem in (P)PDDL:



```
(:action load-box-on-truck-in-city  
:parameters (?b - box ?t - truck ?c - city)  
:precondition (and (Bln ?b ?c) (Tln ?t ?c))  
:effect (prob .9 (and (On ?b ?t) (not (Bln ?b ?c)))))
```

- Can solve *ground MDP* for *each* domain instance:
  - ◆ 3 trucks: 2 planes: 3 boxes:
- Or solve *first-order MDP* for *all* domains at once!
  - ◆ Lift PPDDL problem to first-order MDP (FOMDP)
  - ◆ Solution makes value distinctions for *all* domains!

# Talk Outline

## 1) FOMDP Introduction

- ◆ Original definition, solution (BoutReiPr, IJCAI-01)

## 2) Exploiting structure

- ◆ First-order decision diagrams

## 3) Linear value approximation

## 4) Practical issues & results

## 5) Related work and conclusions

# FOMDP

## Introduction

# Markov Decision Processes

- $\langle S, A, T, R \rangle$ 
  - ◆  $S$ : finite set of states
  - ◆  $A$ : finite set of actions
  - ◆  $T: S \times A \times S \rightarrow [0,1]$  transition function
  - ◆  $R: S \rightarrow \mathbb{R}$  reward function
- Policy  $\pi: S \rightarrow A$
- Value function:  $V(s) = E_{\pi} [\sum_{t=0}^{\infty} \gamma^t r^t | s]$
- $B^a[V(s)] = R(s) + \sum_{s' \in S} T(s, a, s') V(s')$ 
  - ◆  $Q^V(s, a) = B^a[V(s)]$
- $V^*(s) = \max_{a \in A} B^a[V^*(s)]$

# Building Blocks of FOMDPs

- **Case:** Assign value to first-order state abstraction
  - ◆ E.g., express *reward* in *BoxWorld* FOMDP as...

$$rCase(s) = \begin{array}{|l|l|} \hline \forall b,c. Dest(b,c) \Rightarrow BIn(b,c,s) & 1 \\ \hline \neg \text{“} & 0 \\ \hline \end{array}$$

- **Operators:** Define unary, binary case operations
  - ◆ E.g., can take “cross-sum”  $\oplus$  (or  $\otimes$ ,  $\ominus$ ) of cases...

$$\begin{array}{|l|l|} \hline \exists x.A(x) & 10 \\ \hline \neg \exists x.A(x) & 20 \\ \hline \end{array} \oplus \begin{array}{|l|l|} \hline \exists y.A(y) \wedge B(y) & 3 \\ \hline \neg \exists y.A(y) \wedge B(y) & 4 \\ \hline \end{array} = \begin{array}{|l|l|} \hline \exists x.A(x) \wedge \exists y.A(y) \wedge B(y) & 13 \\ \hline \exists x.A(x) \wedge \neg \exists y.A(y) \wedge B(y) & 14 \\ \hline \neg \exists x.A(x) \wedge \exists y.A(y) \wedge B(y) & 23 \\ \hline \neg \exists x.A(x) \wedge \neg \exists y.A(y) \wedge B(y) & 24 \\ \hline \end{array}$$

- ◆ Can remove inconsistent elements, simplify

# FOMDP Foundation: SitCalc

- **Deterministic Actions:**  $\text{loadS}(b,t), \text{unloadS}(b,t), \dots$
- **Fluents (situated):**  $\text{BIn}(b,c,s), \text{TIn}(t,c,s), \text{On}(b,t,s)$
- **Successor-state axioms (SSAs):**
  - ◆ Describe how actions affect fluents
  - ◆ Ex:  $\text{BIn}(b,c, \text{after action } a \text{ in situation } s) \equiv$   
(1) for some  $t$ :  $\text{TIn}(t,c,s) \text{ AND } \text{On}(b,t,s)$   
AND  $a = \text{unloadS}(b,t)$   
OR (2)  $\text{Bin}(b,c,s) \text{ AND } a \neq \text{loadS}(b,t)$
- **Regression Operator:**  $\text{Regr}[\varphi] = \varphi'$ 
  - ◆ Takes a formula  $\varphi$  describing a *post-action* state
  - ◆ Uses SSAs to build  $\varphi'$  describing *pre-action* state

# Stochastic Actions & FODTR

## ■ Stochastic actions using deterministic SitCalc:

- ◆ User's stochastic action:  $A(x) = \text{load}(b,t)$
- ◆ Nature's choice:  $n(x) \in \{\text{loadS}(b,t), \text{loadF}(b,t)\}$
- ◆ Probability distribution over Nature's choice:

$$P(\text{loadS}(b,t) \mid \text{load}(b,t)) = \begin{array}{|c|c|} \hline \text{snow}(s) & .1 \\ \hline \neg \text{snow}(s) & .6 \\ \hline \end{array}$$

$$P(\text{loadF}(b,t) \mid \text{load}(b,t)) = \begin{array}{|c|c|} \hline \text{snow}(s) & .9 \\ \hline \neg \text{snow}(s) & .4 \\ \hline \end{array}$$

## ■ First-order decision-theoretic regression

- ◆ FODTR = *expectation* of regression:

$$\text{FODTR}[\text{vCase}(s), A(x)] = \mathbf{E}_{P(n(x) \mid A(x))} [\text{Regr}[\text{vCase}(s), n(x)] ]$$



# Q-functions and Backups

## ■ FODTR almost gives us a Q-function

$$\text{FODTR}[\text{vCase}(s), \text{unload}(b, t)] = \begin{array}{|c|c|} \hline \text{On}(b, t, s) & 5 \\ \hline \neg \text{On}(b, t, s) & 0 \\ \hline \end{array}$$

- ◆ FODTR specific to action variables
- ◆ Also need to add reward, discount

## ■ Specify a backup operator for this

$$\text{B}^{\text{unload}}[\text{vCase}(s)] = \text{rCase}(s) \oplus \gamma \begin{array}{|c|c|} \hline \exists b, t. \text{On}(b, t, s) & 5 \\ \hline \exists b, t. \neg \text{On}(b, t, s) & 0 \\ \hline \end{array}$$

- ◆ Yields a first-order Q-function


# Symbolic Dynamic Programming

## ■ What value if 0-stages-to-go?

- ◆ Obviously  $V^0(s) = rCase(s)$

## ■ What value if 1-stage-to-go?

- ◆ We know value for each action



- ◆ Now just need max for every state

## ■ Value iteration: (BoutReiPr, IJCAI-01)

- ◆ Obtain  $V^{n+1}$  from  $V^n$  until  $(V^{n+1} \ominus V^n) < \varepsilon$

# Exploiting Structure

# Exploiting Redundancy

- Many case operations in solutions

$\exists x A(x)$	10	$\oplus$	$\exists y A(y)$	1	$=$	$\exists x A(x) \wedge \exists y A(y)$	11
$\neg \exists x A(x)$	20		$\neg \exists y A(y)$	2		<del><math>\exists x A(x) \wedge \neg \exists y A(y)</math></del>	<del>12</del>
						<del><math>\neg \exists x A(x) \wedge \exists y A(y)</math></del>	<del>21</del>
						$\neg \exists x A(x) \wedge \neg \exists y A(y)$	22

- Still have redundant formulae!

- Extract propositional structure

Prop Var	FOL Mapping
a	$\exists x A(x)$
b	$\exists x B(x)$

a	10
$\neg a$	20

 $\oplus$ 

a	1
$\neg a$	2

 $=$ 


a	11
$\neg a$	22

# Exploiting CSI

## ■ First-order ADDs

- ◆ Exploit (FO) context-specific independence

Prop Var	FOL Mapping
a	$\exists x A(x)$
b	$\exists x A(x) \wedge B(x)$



## ■ Some decision paths are unreachable

- ◆ Track implications
- ◆ Avoid inconsistent paths


## ■ Use FOADDs to replace case & operations

# Exploiting Affine Structure

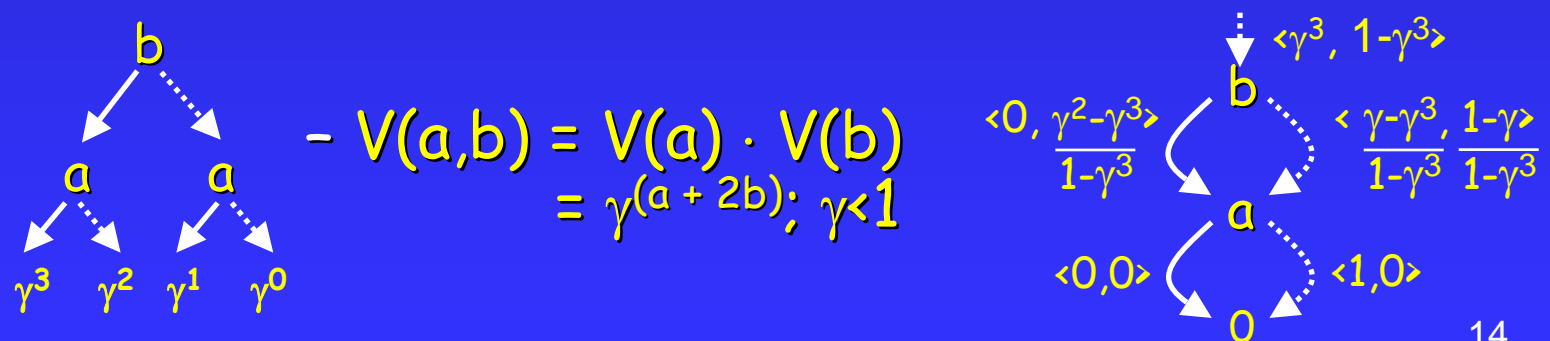
## ■ Replace ADDs with affine ADDs (SanMc,IJCAI-05)

- ◆ Best case: exp→linear reduction; never worse!

## ■ Example 1: Additive reward/utility functions



## ■ Example 2: Multiplicative value functions



# Linear Value Approximation

# First-order Basis Functions

- Approximate value with basis functions:

$$V(s) = w_1 \cdot \begin{array}{|c|c|} \hline \exists b,c \text{ BIn}(b,c,s) & 1 \\ \hline \neg \exists b,c \text{ BIn}(b,c,s) & 0 \\ \hline \end{array} \oplus w_2 \cdot \begin{array}{|c|c|} \hline \exists t,c \text{ TIn}(t,c,s) & 1 \\ \hline \neg \exists t,c \text{ TIn}(t,c,s) & 0 \\ \hline \end{array}$$

- Reduces solution to finding good weights
  - ◆ Weight projection  $\Rightarrow$  no need for simplification
  - ◆ Only need to do consistency checking!
- How to find weights?
  - ◆ Formulate as optimization of LP



# Approximate Linear Programming


- (SanBout, UAI-05) FOALP: Generalize approximate LP solution of van Roy, GKP, SP

◆ Define:  $V(s) = \oplus_{i=1..k} w_i \cdot bCase_i(s)$

Vars:  $w_i; i \leq k$

Minimize:  $\sum_{i=1..k} c_i \cdot w_i$

Subject to:  $V(s) \geq B^a[V(s)]; \forall a \in A, s$



← state space →

# First-order Constraints

Technically  
 $\infty$  constraints

## ■ Example constraint:

$$0 \geq w_1 \cdot \begin{array}{|l|l|} \hline \exists b,c \text{ BIn}(b,c,s) & 1 \\ \hline \neg \exists b,c \text{ BIn}(b,c,s) & 0 \\ \hline \end{array} \oplus w_2 \cdot \begin{array}{|l|l|} \hline \exists t,c \text{ TIn}(t,c,s) & 1 \\ \hline \neg \exists t,c \text{ TIn}(t,c,s) & 0 \\ \hline \end{array} ; \forall s$$

## ■ Only finite *distinct* constraints

$\exists b,c \text{ BIn}(b,c,s) \wedge \exists t,c \text{ TIn}(t,c,s)$	$0 \geq w_1 + w_2$
$\exists b,c \text{ BIn}(b,c,s) \wedge \neg \exists t,c \text{ TIn}(t,c,s)$	$0 \geq w_1$
$\neg \exists b,c \text{ BIn}(b,c,s) \wedge \exists t,c \text{ TIn}(t,c,s)$	$0 \geq w_2$
$\neg \exists b,c \text{ BIn}(b,c,s) \wedge \neg \exists t,c \text{ TIn}(t,c,s)$	$0 \geq 0$

~~$w_1=1, w_2=1$~~

~~$w_1=-1, w_2=1$~~

$w_1=-1, w_2=-1$




## ■ Solve via constraint generation

- ◆ Efficiently find max violated constraints
- ◆ Generalize variable elimination to first-order

# Policy Construction

- Derive greedy policy  $\pi$  from  $V(s)$ :



- Now build  $\pi\text{Case}_a(s)$  for  $a \in \{\text{unload}, \text{load}\}$ :

$\pi\text{Case}_{\text{unload}}(s) =$


$\exists b,t \varphi_1$	9
$\neg \exists b,t \varphi_1 \wedge \neg \exists b,t \varphi_3$ $\wedge \neg \exists b,t \varphi_4 \wedge \exists b,t \varphi_2$	0

$\pi\text{Case}_{\text{load}}(s) =$

$\neg \exists b,t \varphi_1 \wedge \exists b,t \varphi_3$	3
$\neg \exists b,t \varphi_1 \wedge \neg \exists b,t \varphi_3$ $\wedge \exists b,t \varphi_4$	1

# Approximate Policy Iteration

## ■ Basic Algorithm

- 1) Define  $V^{(j)}(s) = \oplus_{i=1..k} w_i^{(j)} \cdot bCase_i(s)$ ,  
Initialize  $j=0$ ,  $\pi=\{\text{any policy}\}$
  - 2) Given policy  $\pi^j$ , find Bellman-error  
minimizing  $w_i^{(j)}$  for  $V^{(j)}(s)$
  - 3) Derive greedy policy  $\pi^{j+1}$  from  $V^{(j)}(s)$
  - 4) If  $\pi^{j+1} \neq \pi^j$  then let  $j=j+1$  and go to step 2
- 

## ■ LP for Bellman-error minimizing $w_i^{(j)}$ :

Vars:  $w_i^{(j)}; i \leq k$

Minimize:  $\phi$

Subject to:  $\phi \geq | \pi Case_a^{(j)}(s) \oplus V^{(j)}(s) \ominus B^a[V^{(j)}(s)] |; \forall a \in A, s$

## ■ Use $\pi Case_a(s)$ to enforce $B^\pi$ (GKPV, JAIR-02)

# Practical Issues & Results

# Generating Basis Functions

## ■ Where do basis functions come from?

- ◆ Major question for automation
- ◆ Systematically build from FOL components?
- ◆ Candidate space too large!

## ■ Idea (Gretton & Thiebaux, UAI-04) :

- ◆ Regressions from goal make good candidates
- ◆ Guaranteed to have some value
- ◆ Building blocks of value iteration

## ■ Iteratively solve FOMDP

- ◆ Retain basis functions with weight  $>$  threshold
- ◆ Generate new basis functions from retained set

# Problems w/ Universal Reward

- Universal rewards are difficult for FOMDPs, e.g.,

- ◆ Given reward:

$$r_{\text{Case}}(s) = \begin{array}{|l|l|} \hline \forall b, c. \text{Dest}(b, c) \Rightarrow \text{BIn}(b, c, s) & 1 \\ \hline \neg & 0 \\ \hline \end{array}$$

- ◆ Exact n-stage-to-go value function has form:

$$v_{\text{Case}}^n(s) = \begin{array}{|l|l|} \hline \forall b, c. \text{Dest}(b, c) \Rightarrow \text{BIn}(b, c, s) & 1 \\ \hline 1 \text{ box not at dest} & \gamma \\ \hline \dots & \dots \\ \hline n-1 \text{ boxes not at dest} & \gamma^{n-1} \\ \hline \end{array}$$

- ◆ Exact value function has infinitely many values!
- ◆ No compact representation  
(using piecewise-constant case statement)

# Additive Goal Decomposition

## ■ Off-line solution for universal rewards:

- ◆ Given goal  $\forall b,c. \text{Dest}(b,c) \Rightarrow \text{BIn}(b,c,s)$
- ◆ Solve FOMDP for goal  $\text{BIn}(b^*,c^*,s)$  to get  $V(b^*,c^*,s)$

## ■ At run-time:

- ◆ Given concrete domain:  $\{\text{Dest}(b_1,c_1), \text{Dest}(b_2,c_2)\}$
- ◆ “Score” actions additively w.r.t. each goal


$$\begin{aligned}
 & \text{B}^{\text{unload}}[V(b_1,c_1,s)] \quad \oplus \quad \text{B}^{\text{unload}}[V(b_2,c_2,s)] \\
 &= \begin{array}{|c|c|} \hline \exists b,t \ \varphi_1 & 3 \\ \hline \exists b,t \ \varphi_2 & 1 \\ \hline \end{array} \quad \oplus \quad \begin{array}{|c|c|} \hline \exists b,t \ \varphi_3 & 3 \\ \hline \exists b,t \ \varphi_4 & 1 \\ \hline \end{array} \\
 & \{ \text{unload}(b_1,t_1) \rightarrow 3, \text{unload}(b_2,t_2) \rightarrow 1 \} \quad + \quad \{ \text{unload}(b_2,t_2) \rightarrow 3, \text{unload}(b_2,t_1) \rightarrow 1 \} = \boxed{\begin{array}{l} \text{unload}(b_1,t_1) \rightarrow 3, \\ \text{unload}(b_2,t_2) \rightarrow 4, \\ \text{unload}(b_2,t_1) \rightarrow 1 \end{array}}
 \end{aligned}$$




# Optimizations

- **Enforce disjointness in basis functions:**
  - ◆ Can reduce search of  $2^{|B|}$  case partitions to  $|B|$
- **Exploiting implicit max in constraint generation**
  - ◆ Don't always need to enforce disjointness
  - ◆ Max-sum does this automatically
- **FOADDs for formulae simplification**
- **Huge cache of proved/unproved theorems**
  - ◆ Store FOL formulae in canonical format
- **Structural optimization in CNF transformation**
  - ◆ Introduce propositional literals to exploit DPLL in Vampire
- **Join-order optimizations in policy matcher**


# Results – ICAPS-06, Elevators



# Results – ICAPS-06, Tire World



# Results – ICAPS-06, Blocks World



# Related Work

- Direct “first-order” value iteration:
  - ◆ ReBel algorithm for RMDPs (KvOdR, 2004)
  - ◆ FOVIA algorithm for fluent calculus (KS, 2005)
  - ◆ First-order decision diagrams (JKW, 2007)
  - ◆ → all disallow  $\forall$  quant., e.g., universal cond. effects
- Sampling and/or inductive techniques:
  - ◆ Approx. linear programming for RMDPs (GKGK, 2003)
  - ◆ Inductive policy selection using FO regression (GT, 2004)
  - ◆ Approximate policy iteration (FYG, 2004)
  - ◆ → sampled domain instantiations do not ensure generalization across all possible worlds
  - ◆ → must restrict to “small” domain instances

# Conclusions and Future Work

## ■ Conclusions:

- ◆ Managing structure in FOMDP solutions
- ◆ Approximation techniques for FOMDPs
- ◆ Range of practical implementation issues
- ◆ **Only completely first-order planner to date**
  - ◆  $\Rightarrow$  2<sup>nd</sup> place in ICAPS 2006 IPPC by # problems solved

## ■ Current & future work:

- ◆ Sum aggregator:  $\sum_c \exists c \text{ BIn}(b, c, s): 1$ ; factored actions
- ◆ Program constraints
- ◆ Handling real-valued quantities, arithmetic
- ◆ Exploiting topological structure
- ◆ Integration with RL? First-order POMDPs?

# Bibliography

- (BRP, 2001) C. Boutilier, R. Reiter, and B. Price. (2001). *Symbolic Dynamic Programming for First-order MDPs*. IJCAI-01.
- (SP, 2001) Dale Schuurmans and Relu Patrascu. (2001) Direct value approximation for factored MDPs. NIPS-2001.
- (SM, 2005) S. Sanner and D. McAllester. (2005). *Affine Algebraic Decision Diagrams (AADDs) and their Application to Structured Probabilistic Inference*. IJCAI-05.
- (SB, 2005) S. Sanner and C. Boutilier. (2005). *Approximate Linear Programming for First-order MDPs*. UAI-05.
- (GKGK, 2003) C. Guestrin, D. Koller, C. Gearhart, and N. Kanodia. (2003). Generalizing Plans to New Environments in Relational MDPs. IJCAI-03.
- (GT, 2004) C. Gretton and S. Thiebaux. (2004). Exploiting first-order regression in inductive policy selection. UAI-04.
- (GKPV, 2002) C. Guestrin, D. Koller, R. Parr, and S. Venktaraman. Efficient solution methods for factored MDPs. JAIR, 2002.

# Extra: Language Extensions



# Sum & Product Aggregators

- Often, reward scales with domain size:

$$rCase(s) = \sum_c \begin{array}{|l|l|} \hline \exists c \text{ BIn}(b,c,s) & 1 \\ \hline \neg \exists c \text{ BIn}(b,c,s) & 0 \\ \hline \end{array}$$

- Beyond expressive power of current FOMDP
- Need language extension for sum/product aggregators
  - ◆ Functional as opposed to truth semantics
  - ◆ Like a quantifier (but indefinite  $\oplus$ )
- Can extend symbolic dynamic programming, approximate solutions
  - ◆ But tricky

# Factored Actions

- What if action has indefinite number of independent outcomes?

$$P(\text{lost}(b) \mid a) =$$

large(b)	.0001
medium(b)	.0005
small(b)	.001

- Then we get an indefinitely large joint distribution:

$$P(\text{lost}(b_1) \circ \dots \circ \text{lost}(b_n) \mid a) = \prod_b$$

large(b)	.0001
medium(b)	.0005
small(b)	.001

- Have to exploit (FO) independence in solutions
  - ◆ Then most of product will marginalize